

《并行计算》上机报告

姓名:	范文	学号:	PB18111679	日期:	2021.5.22
上机题目:	MPI 并行编程实验				
实验环境:					
CPU: 8 个 Intel(R) Core(TM) i5-8250U					
内存: 8G (交换区 8G)					
操作系统: ubuntu 20.04 (内核版本为 5.4.0-72-generic)					
软件平台: VSCode,					
一、算法设计与分析:					
题目一:					
根据 PPT 上的内容, 可以在求和近似积分的基础上, 设计基于消息传递的并行求 π 的算法:					
(1) 非 0 号线程执行局部的求和, 并将其计算结果传递给 0 号线程。					
(2) 0 号线程将这些结果进行累加, 即得到 π 的近似值。					
相比于 OpenMP 中共享存储的方法进行局部结果的累加, mpi 通过消息传递的方式需要考虑更多的开销, 因此更慢一些。					

题目二：

根据教材上的内容，可以把 PSRS 分为以下 8 个阶段进行：

- (1) 均匀划分: 将 n 个元素 $A[1..n]$ 均匀划分成 p 段,
- (2) 局部排序: p_i 调用串行排序算法对 $A[(i-1)n/p+1..in/p]$ 排序
- (3) 选取样本: p_i 从其有序子序列 $A[(i-1)n/p+1..in/p]$ 中选取 p 个样本元素
- (4) 样本排序: 用一台处理器对 p^2 个样本元素进行串行排序
- (5) 选择主元: 用一台处理器从排好序的样本序列中选取 $p-1$ 个主元, 并播送给其他 p_i
- (6) 主元划分: p_i 按主元将有序段 $A[(i-1)n/p+1..in/p]$ 划分成 p 段
- (7) 全局交换: 各处理器将其有序段按段号交换到对应的处理器中
- (8) 归并排序: 各处理器对接收到的元素进行归并排序

分析以上过程可知, (1)(4)(5)(7)(8)可以在一个线程上执行, (2)(3)(6)可以多个线程并行执行。而不同步骤之间需要进行一到多或者多到一的消息传递。

如果不考虑通信的开销, 时间复杂度的分析和 OpenMP 的类似:

- (1) 均匀划分需要 $O(p)$ 的复杂度。
- (2) 局部排序需要 $O(n/p \lg(n/p))$ 的复杂度。
- (3) 选取样本需要 $O(p)$ 的复杂度。
- (4) 样本排序需要 $O(p^2 \lg p)$ 的复杂度。
- (5) 选择主元需要 $O(p)$ 的复杂度。
- (6) 主元划分需要 $O(p)$ 的复杂度。

(7) 全局交换需要 $O(p^2)$ 的复杂度。

(8) 归并排序需要 $O(n/p \lg(n/p))$ 的复杂度。

因此，不考虑通信开销，基于 mpi 的 PSRS 算法的复杂度为 $O(n/p \lg(n/p) + p^2 \lg p)$ ；但是，通信开销会使这里的复杂度继续增加。

二、核心代码：

题目一：

(1) 在 main 函数中，控制不同线程号的线程执行不同的操作。

```
// for ID != 0, calculate the partial sum over step
if(thread_ID != 0){
    calculate_over_step(thread_ID,thread_num-1);
}
// for ID = 0, aggregate the partial sums
else{
    double pi = aggregate(thread_num);
    std::cout << "pi is " << std::setprecision(8) << pi << std::endl;
}
```

(2) 在非 0 号线程执行的 calculate_over_step() 函数中，先进行循环的局部求和，再把求和的结果通过 MPI_Send() 发给 0 号线程。

```
// calculate the integral
// by getting sum of the interval over step
// @_thread_ID: current thread ID
// @_thread_num: total number of threads for computing integral
// send back value of the partial integral for aggregation
void calculate_over_step(int _thread_ID, int _thread_num){
    double x;
    double step = 1.0 / num_steps;
    double partial_sum = 0.0;
    // calculate the partial sum over step
    for(int i = _thread_ID; i < num_steps; i += _thread_num){
        x = (i + 0.5)*step;
        partial_sum += 4.0/(1.0 + x*x);
    }
    partial_sum *= step;
    // send the partial_sum to thread 0
    MPI_Send(&partial_sum, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD);
}
```


(3) 在 0 号线程执行的 aggregate()函数中,通过 MPI_Recv() 接收每个非 0 线程的求和的结果, 并将这些局部的结果进行累加。

```
38 // aggregate the partial sum computed by thread 1 ~ _thread_num
39 // @_thread_num: the number of threads
40 double aggregate(int _thread_num){
41     double pi = 0.0;
42     double partial_sum;
43     // receive the partial_sum calculated by other threads
44     // and aggregate them
45     for(int i = 0; i < _thread_num - 1; ++i){
46         MPI_Recv(&partial_sum, 1, MPI_DOUBLE, MPI_ANY_SOURCE,
47                 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
48         pi += partial_sum;
49     }
50     return pi;
51 }
```

题目二:

PSRS 算法的例子如下所示

■ 例6.1 PSRS排序过程。N=27, p=3, PSRS排序如下:



由此可知, (1)均匀划分、(4)采样排序、(5)选择主元、(7)全局交换、(8)归并排序都是在同一个线程里完成的。而(2)局部排序、(3)正则采样、(6)主元划分都是多线程并行完成

的。因此，我令 0 号线程执行 serial() 串行函数，令非 0 号线程执行 parallel() 并行函数，完成以上操作；而核心代码如下：

一. 0 号线程读取输入。

```
57 // get the input from the file
58 int elem_num;
59 int* array = get_input(elem_num);
60 // show the input array
61 std::cout << "original array is\n";
62 for(int i = 0; i < elem_num; ++i){
63     std::cout << array[i] << " ";
64 }
65 std::cout << "\n";
```

二. 0 号线程对输入元素进

```
67 // -----
68 // (1) uniform partition
69 // -----
70 // the index of start and end for each partition
71 int start_index, end_index;
72 // the gap between two neighbor partition
73 int gap = (int)ceil(elem_num / _thread_num);
74 for(int i = 1; i <= _thread_num; ++i){
75     start_index = gap * (i - 1);
76     end_index = gap * i < elem_num ? gap * i : elem_num;
77     int partition_size = end_index - start_index;
78     // send the elements of each partition
79     // to the corresponding thread
80     MPI_Send(&partition_size, 1, MPI_INT, i, 1, MPI_COMM_WORLD);
81     MPI_Send(array + start_index,
82             end_index - start_index, MPI_INT, i, 2, MPI_COMM_WORLD);
83 }
```

三. 每个非 0 号线程都得到它们对应的划分，并对它们进行排序。


```

163 // receive the uniform partition from thread 0
164 int array_size;
165 MPI_Recv(&array_size,1, MPI_INT, 0, 1,MPI_COMM_WORLD, MPI_STATUS_IGNORE);
166 int* local_array = new int[array_size];
167 MPI_Recv(local_array, array_size, MPI_INT, 0, 2,
168         MPI_COMM_WORLD, MPI_STATUS_IGNORE);
169
170 // -----
171 // (2) local sort
172 // -----
173 std::sort(local_array,local_array + array_size);
174

```

四. 每一个非 0 号线程对划分进行采样，并传递给 0 号线程。

```

175 // -----
176 // (3) normal sampling
177 // -----
178 // get _thread_num samples in a partition
179 int gap = (int)floor( array_size / _thread_num );
180 int* local_samples = new int[_thread_num];
181 int index = 0;
182 for(int i = 0;i < _thread_num;++i){
183     local_samples[i] = local_array[index];
184     index += gap;
185 }
186 // send the samples to thread_0 for next step
187 MPI_Send(local_samples,_thread_num, MPI_INT, 0, 3, MPI_COMM_WORLD);
188

```

五. 0 号线程收集采样，并将采样进行排序

```

85 // receive the samples from those threads
86 int* samples = new int [_thread_num * _thread_num];
87 // the offset in samples for each reception
88 int offset = 0;
89 for(int i = 0;i < _thread_num;++i){
90     MPI_Recv(samples + offset, _thread_num,
91             MPI_INT,MPI_ANY_SOURCE, 3, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
92     offset += _thread_num;
93 }
94
95 // -----
96 // (4) sample sorting
97 // -----
98 std::sort(samples,samples + offset);

```

六. 0 号线程选择主元, 并将主元传递给每个非 0 号线程。

```
100 // -----
101 // (5) pivot select
102 // -----
103 int *pivots = new int[_thread_num - 1];
104 // get the pivots uniformly
105 // the step is _thread_num
106 for(int i = 0; i < _thread_num - 1; ++i){
107     pivots[i] = samples[ (i + 1) * _thread_num ];
108 }
109 // send the pivots to thread 1 ~ _thread_num
110 for(int i = 1; i <= _thread_num; ++i){
111     MPI_Send(pivots, _thread_num - 1, MPI_INT, i, 4, MPI_COMM_WORLD);
112 }
113
```

七. 每一个非 0 号线程接收主元, 并根据主元再次进行划分, 用一个二维数组存放, 并将这些划分传递给 0 号线程。

```
189 // -----
190 // (6) privot partition
191 // -----
192 // receive the pivots sent from the first partition
193 int *local_pivots = new int[_thread_num - 1];
194 MPI_Recv(local_pivots, _thread_num - 1, MPI_INT,
195         0, 4, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
196
197 // partition according to the pivots
198 // replace_list[i] means the elements to be sent to thread i
199 std::array<std::vector<int>, MAX_THREAD_NUM> replace_list;
200 int pivot_index = 0;
201 for(int i = 0; i < array_size;){
202     if(local_array[i] < local_pivots[pivot_index]){
203         replace_list[pivot_index].push_back(local_array[i]);
204         i++;
205     }
206     else if(pivot_index == _thread_num - 1){
207         replace_list[pivot_index - 1].push_back(local_array[i]);
208         i++;
209     }
210     else{
211         pivot_index++;
212     }
213 }
```

八. 0 号线程接收这些进一步的划分, 并对他们进行归并排序, 得到了最终有序的数组。

```
114 // -----
115 // (7) global exchange & (8) merge sort
116 // -----
117 // receive the partitions sent by thread 1 ~ _thread_num - 1
118 std::array<std::vector<int>, MAX_THREAD_NUM> replace_list;
119 for (int thread_index = 1; thread_index <= _thread_num; ++ thread_index){
120     for(int part_index = 0; part_index < _thread_num; ++ part_index){
121         // get the partition in a thread
122         unsigned len;
123         MPI_Recv(&len, 1, MPI_UNSIGNED, thread_index, 5, MPI_COMM_WORLD,
124             MPI_STATUS_IGNORE);
125         std::vector<int> partition(len, 0);
126         MPI_Recv(&partition[0], len, MPI_INT, thread_index, 6,
127             MPI_COMM_WORLD, MPI_STATUS_IGNORE);
128         int next_part_index = partition[len-1];
129         // tmp for merging
130         std::vector<int> tmp;
131         std::merge(replace_list[next_part_index].begin(), replace_list
132             [next_part_index].end(),
133             partition.begin(), --partition.end(), std::back_inserter(tmp));
134         replace_list[next_part_index] = tmp;
135     }
136 }
```

三、结果与分析:

题目一:

我编译了程序之后, 使用 3 个线程进行执行 (实际上只有两个线程进行局部求和), 得到了结果如下:

```
eddie@eddie-TM1701:~/Desktop/junior_2/parallel/LAB/lab2$ mpicxx -o par_pi par_pi.cpp
eddie@eddie-TM1701:~/Desktop/junior_2/parallel/LAB/lab2$ mpirun -np 3 ./par_pi
pi is 3.1415527
```

这里的结果精确到了小数点后 4 位。

题目二：

使用 PPT 上的例子，开 4 个线程执行，得到测试结果如下：

```
eddie@eddie-TM1701:~/Desktop/junior_2/parallel/LAB/lab2$ mpicxx -o PSRS PSRS.cpp
eddie@eddie-TM1701:~/Desktop/junior_2/parallel/LAB/lab2$ mpirun -np 4 ./PSRS
original array is
15 46 48 93 39 6 72 91 14 36 69 40 89 61 97 12 21 54 53 97 84 58 32 27 33 72 20
after sorting, the array is
6 12 14 15 20 21 27 32 33 36 39 40 46 48 53 54 58 61 69 72 72 84 89 91 93 97 97
```

这说明可以正确排序。

四、备注（* 可选）：

有可能影响结论的因素：

1. 计算精度会导致估计 π 的结果并不是很准确
2. mpi 中的线程的消息传递和并行调度会产生开销，导致计算并不是很快。（也可能是输入的数据量较小，没有明显体现出并行编程的优越性出来）

总结：

本次实验，我学习了使用 mpi 进行并行编程，了解了基于消息传递的并行编程的思想，并提高了自己的编码和调试能力。

附录（源代

算法源代码（C/C++/JAVA 描述）

码)

我的本次实验的源码保存在了 https://github.com/fanweneddie/parallel_computing/tree/master/lab2 中。

题目 1:

```
// par_pi.c
/*
 * using MPI to calculate the approximate value of Pi
 * by computing parallelism integral over step
 */
/* I have borrowed some ideas from myl7's github
 * https://github.com/myl7/paracomp2021
 */
/* PB18111679 fanweneddie
 */
#include <iostream>
#include <iomanip>
#include <mpi.h>

static long num_steps = 100000;

void calculate_over_step(int_thread_ID,int_thread_num);
double aggregate(int_thread_num);

// calculate the integral
// by getting sum of the interval over step
// @_thread_ID: current thread ID
// @_thread_num: total number of threads for computing integral
// send back value of the partial integral for aggregation
void calculate_over_step(int_thread_ID,int_thread_num){
double x;
double step = 1.0 / num_steps;
double partial_sum = 0.0;
// calculate the partial sum over step
for(int i = _thread_ID;i < num_steps;i += _thread_num){
x = (i + 0.5)*step;
partial_sum += 4.0/(1.0 + x*x);
}
partial_sum *= step;
// send the partial_sum to thread 0
```

```
MPI_Send(&partial_sum,1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD);
}

// aggregate the partial sum computed by thread 1 ~ _thread_num - 1
// @ _thread_num: the number of threads
double aggregate(int _thread_num){
double pi = 0.0;
double partial_sum;
// receive the partial_sum calculated by other threads
// and aggregate them
for(int i = 0; i < _thread_num - 1; ++i){
MPI_Recv(&partial_sum,1,MPI_DOUBLE,MPI_ANY_SOURCE,
0,MPI_COMM_WORLD,MPI_STATUS_IGNORE);
pi += partial_sum;
}
return pi;
}

int main(int argc, const char *const argv[]) {

MPI_Init(&argc, const_cast<char ***>(&argv));
// the thread ID of the current thread
int thread_ID;
MPI_Comm_rank(MPI_COMM_WORLD, &thread_ID);
// the total number of threads
int thread_num;
MPI_Comm_size(MPI_COMM_WORLD, &thread_num);

// for ID != 0, calculate the partial sum over step
if(thread_ID != 0){
calculate_over_step(thread_ID,thread_num-1);
}

// for ID = 0, aggregate the partial sums
else{
double pi = aggregate(thread_num);
std::cout << "pi is " << std::setprecision(8) << pi << std::endl;
}

MPI_Finalize();
return 0;
}
```

题目二：

```
/*
 * use mpi to implement PSRS sorting algorithm
 *
 * the process of PSRS
 * comes from our course slide
 *
 * I have borrowed some ideas from myl7's github
 * https://github.com/myl7/paracomp2021
 *
 * PB18111679 fanweneddie
 */

#include <fstream>
#include <iostream>
#include <array>
#include <vector>
#include <math.h>
#include <algorithm>
#include <mpi.h>

constexpr auto MAX_THREAD_NUM = 10;
void serial(int _thread_num);
void parallel(int _thread_ID, int _thread_num);
int* get_input();

// get the input int array from the input file
// allocate the array and return it
// @_elem_num: the number of elements in the array
// we need to revise it in the function
int* get_input(int& _elem_num){
    std::ios::sync_with_stdio(false);
    std::ifstream in("input.txt");
    if (!in) {
        std::cerr << "Getting input failed\n";
    }
    in >> _elem_num;
    int *array = new int[_elem_num];
    for (int i = 0; i < _elem_num; i++) {
        in >> array[i];
    }
    in.close();
}
```

```

return array;
}

// serial working part
// for thread 0
// do the serial work in PSRS,
// that is, to do
// (1) uniform partition
// (4) sample sorting
// (5) pivot select
// (7) global exchange
// (8) merge sort
// @_thread_num: number of parallel threads
void serial(int _thread_num){

// get the input from the file
int elem_num;
int* array = get_input(elem_num);
// show the input array
std::cout << "original array is\n";
for(int i = 0; i < elem_num; ++i){
std::cout << array[i] << " ";
}

std::cout << "\n";

// -----
// (1) uniform partition
// -----
// the index of start and end for each partition
int start_index, end_index;
// the gap between two neighbor partition
int gap = (int)ceil(elem_num / _thread_num);
for(int i = 1; i <= _thread_num; ++i){
start_index = gap * (i - 1);
end_index = gap * i < elem_num ? gap * i : elem_num;
int partition_size = end_index - start_index;
// send the elements of each partition
// to the corresponding thread
MPI_Send(&partition_size, 1, MPI_INT, i, 1, MPI_COMM_WORLD);
MPI_Send(array + start_index,
end_index - start_index, MPI_INT, i, 2, MPI_COMM_WORLD);
}

```

```

// receive the samples from those threads
int* samples = new int[_thread_num * _thread_num];
// the offset in samples for each reception
int offset = 0;
for(int i = 0; i < _thread_num; ++i){
MPI_Recv(samples + offset, _thread_num,
MPI_INT, MPI_ANY_SOURCE, 3, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
offset += _thread_num;
}

// -----
// (4) sample sorting
// -----
std::sort(samples, samples + offset);

// -----
// (5) pivot select
// -----
int* pivots = new int[_thread_num - 1];
// get the pivots uniformly
// the step is _thread_num
for(int i = 0; i < _thread_num - 1; ++i){
pivots[i] = samples[(i + 1) * _thread_num];
}

// send the pivots to thread 1 ~ _thread_num
for(int i = 1; i <= _thread_num; ++i){
MPI_Send(pivots, _thread_num - 1, MPI_INT, i, 4, MPI_COMM_WORLD);
}

// -----
// (7) global exchange & (8) merge sort
// -----
// receive the partitions sent by thread 1 ~ _thread_num - 1
std::array<std::vector<int>, MAX_THREAD_NUM> replace_list;
for(int thread_index = 1; thread_index <= _thread_num; ++thread_index){
for(int part_index = 0; part_index < _thread_num; ++part_index){
// get the partition in a thread
unsigned len;
MPI_Recv(&len, 1, MPI_UNSIGNED, thread_index, 5, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
std::vector<int> partition(len, 0);
MPI_Recv(&partition[0], len, MPI_INT, thread_index, 6, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
int next_part_index = partition[len-1];
// tmp for merging

```

```

std::vector<int> tmp;
std::merge(replace_list[next_part_index].begin(), replace_list[next_part_index].end(),
partition.begin(), --partition.end(), std::back_inserter(tmp));
replace_list[next_part_index] = tmp;
}
}

// new array to store the sorted elements
std::vector<int> new_array;
// concat the replace_lists to get the new_array
for(int i = 0; i < _thread_num; ++i){
new_array.insert(new_array.end(), replace_list[i].begin(), replace_list[i].end());
}

std::cout << "after sorting, the array is\n";
for (auto elem : new_array) {
std::cout << elem << " ";
}
std::cout << std::endl;

delete [] array;
delete [] samples;
delete [] pivots;
}

// parallel working part
// for thread 1 ~ _thread_num - 1
// do the parallel work in PSRS
// that is, to do
// (2) local sort
// (3) uniform sampling
// (6) pivot partition
void parallel(int _thread_ID, int _thread_num){
// receive the uniform partition from thread 0
int array_size;
MPI_Recv(&array_size, 1, MPI_INT, 0, 1, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
int* local_array = new int[array_size];
MPI_Recv(local_array, array_size, MPI_INT, 0, 2,
MPI_COMM_WORLD, MPI_STATUS_IGNORE);

// -----
// (2) local sort
// -----
std::sort(local_array, local_array + array_size);

```

```

// -----
// (3) normal sampling
// -----
// get _thread_num samples in a partition
int gap = (int)floor( array_size / _thread_num );
int* local_samples = new int[_thread_num];
int index = 0;
for(int i = 0; i < _thread_num; ++i){
    local_samples[i] = local_array[index];
    index += gap;
}

// send the samples to thread_0 for next step
MPI_Send(local_samples, _thread_num, MPI_INT, 0, 3, MPI_COMM_WORLD);

// -----
// (6) pivot partition
// -----
// receive the pivots sent from the first partition
int *local_pivots = new int[_thread_num - 1];
MPI_Recv(local_pivots, _thread_num - 1, MPI_INT,
0, 4, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
// partition according to the pivots
// replace_list[i] means the elements to be sent to thread i
std::array<std::vector<int>, MAX_THREAD_NUM> replace_list;
int pivot_index = 0;
for(int i = 0; i < array_size;){
    if(local_array[i] < local_pivots[pivot_index]){
        replace_list[pivot_index].push_back(local_array[i]);
        i++;
    }
    else if(pivot_index == _thread_num - 1){
        replace_list[pivot_index - 1].push_back(local_array[i]);
        i++;
    }
    else{
        pivot_index++;
    }
}

// send the result of partition to thread 0
for(int i = 0; i < _thread_num; ++i){
    int len = (unsigned)replace_list[i].size();
    // marks the object of this partition.

```

```
replace_list[i].push_back(i);
len++;
MPI_Send(&len, 1, MPI_UNSIGNED, 0, 5, MPI_COMM_WORLD);
MPI_Send(&replace_list[i][0], len, MPI_INT, 0, 6, MPI_COMM_WORLD);
}

delete [] local_array;
delete [] local_samples;
delete [] local_pivots;
}

int main(int argc, const char *const argv[]) {
MPI_Init(&argc, const_cast<char ***>(&argv));
// the thread ID of the current thread
int thread_ID;
MPI_Comm_rank(MPI_COMM_WORLD, &thread_ID);
// the total number of threads
int thread_num;
MPI_Comm_size(MPI_COMM_WORLD, &thread_num);
thread_num = 4;
// for ID = 0, do the serial work
if(thread_ID == 0){
serial(thread_num - 1);
}
// for ID != 0, do parallel work
else{
parallel(thread_ID,thread_num - 1);
}
MPI_Finalize();
return 0;
}
```